



UNIVERSITY OF CARTHAGE
NATIONAL INSTITUTE OF APPLIED SCIENCES AND
TECHNOLOGY

GraphQL and Hasura

Author
Mohamed Helmi BOUDHRAA

December 8, 2020

Contents

1 GraphQL's Scheme Definition Language (SDL)	6
1.1 Definition	6
1.2 Types	6
1.2.1 Object Types	6
1.2.2 Built-in Scalar Types	7
1.3 Enum	7
1.4 Interface	7
1.5 Scheme Directive	8
2 Queries, Mutations and Subscriptions	9
2.1 Introduction	9
2.2 Queries	9
2.2.1 Get all users	9
2.2.2 Get user by id	10
2.2.3 Get user by field name	10
2.2.4 Get users with fragments	10
2.2.5 Use variables	10
2.3 Mutations	11
2.4 Subscriptions	11
3 GraphQL VS. REST	12
3.1 Introduction	12
3.2 Comparison Table	12
3.2.1 Overfetching	12
3.2.2 Underfetching	12
3.2.3 Single Endpoint VS. Multiple Endpoints	13
4 Practical case: Hasura	14
4.1 Introduction	14
4.2 Steps	14
4.2.1 Create a new docker-compose file locally	14
4.2.2 Copy the docker-compose content from GraphQL Github Repository	14
4.2.3 Check the local file variables	15
4.2.4 Run the docker-compose up command	15
4.2.5 Open the console on port 8080	16
4.2.6 Create a new table: Users	16
4.2.7 Create a new table: Courses	17
4.2.8 Configure Primary and Foreign Keys in Courses	17
4.2.9 Define the relationship between tables	18

4.2.10	Mutation: Adding new users	18
4.2.11	Mutation: Adding new courses	19
4.2.12	Subscription: Subscribing to users	19

List of Figures

1.1	User Type	6
1.2	Presentation of object types	7
1.3	Gender enum	7
1.4	Interface Definition	8
1.5	Applying directive to name	8
2.1	Get the names of all the users	9
2.2	Get name of the user of id 1	10
2.3	Get user ages with name "iyadh"	10
2.4	Get name and age of all users	10
2.5	Set id as a variable that can be an input by the client	10
2.6	Creating a new user	11
3.1	REST Example Endpoint	13
3.2	REST Example Endpoint 2	13
3.3	GraphQL Example Endpoint	13
4.1	File creating command	14
4.2	Github File	14
4.3	Docker-compose local file	15
4.4	Command Running	15
4.5	GraphQL Console	16
4.6	Users Table Creation	16
4.7	Courses Table Creation	17
4.8	Key configuration	17
4.9	Relationship Definition	18
4.10	Adding users	18
4.11	Adding courses	19
4.12	User subscription	19

Introduction

In 2012, Facebook internally developed a new Query Language for APIs and a server-side runtime for executing queries called GraphQL, which was publicly released in 2015 to join the open-source community. Later in 2018, a new foundation, GraphQL Foundation, was established to adopt the project from Facebook, aiming to make it omnipresent across web platforms.

GraphQL's approach to developing APIs is comparable and different in a certain way from other web service architectures such as REST. It enables declarative data fetching, where the client can specify what data he needs from an API, without going through multiple endpoints. And all that by defining a certain structure and a single endpoint of communication.

In this report, we will mainly cover different aspects of GraphQL as follows:

First, we will talk about Schema Definition Language of GraphQL as well as its different components.

Second of all, we will go through CRUD (Create Read Update Delete) and subscription using GraphQL's queries, mutations and subscriptions.

After that, we will illustrate a high-level comparison between GraphQL and REST to better understand the pros of migrating.

Finally, we will go through a practical example using Hasura GraphQL Engine.

Chapter 1

GraphQL's Scheme Definition Language (SDL)

1.1 Definition

A GraphQL Schema Definition is the most concise way to specify a GraphQL schema. The syntax takes part in the official GraphQL specification as a well-defined structure. [Bur17] The GraphQL schema for a blogging app could be specified like this:

```
type User {  
  id: Int!  
  name: String!  
  age: Int!  
}
```

Figure 1.1: User Type

1.2 Types

1.2.1 Object Types

A GraphQL service has specific object types, which are defined with the **type** keyword and start with a capital letter by convention. [All18]

Object types can be the following:

- Query
- Mutation
- Subscription

```

type User {
  id: Int!
  name: String!
  age: Int!
}

type Query {
  user(id: Int!): User!
  allUsers: [User!]!
}

type Mutation {
  addUser(name: String!, age: Int!): User!
  removeUser(id: Int!): User!
}

schema {
  query: Query
  mutation: Mutation
}

```

Figure 1.2: Presentation of object types

1.2.2 Built-in Scalar Types

GraphQL has 5 built-in scalar types, which are the following:

- Int
- Float
- String
- Boolean
- ID (resolves to a string, but expects a unique value)

1.3 Enum

An enum is a scalar value that has a specified set of possible values. In the example down below, we illustrate a Gender enum to limit the specifications for the user.

```

type User {
  id: Int!
  name: String!
  age: Int!
  gender: Gender!
}

enum Gender {
  Female
  Male
}

```

Figure 1.3: Gender enum

1.4 Interface

An interface is defined as a list of fields. A GraphQL type must have the same fields as all the interfaces it implements and all interface fields must be of the same type. In the example down below, we present a User interface to be inherited by multiple types of users.

```

interface User {
  id: Int!
  name: String!
  age: Int!
  gender: Gender!
}

type Student implements User {
  id: Int!
  name: String!
  age: Int!
  gender: Gender!
  level: Int!
  major: String!
}

enum Gender {
  Female
  Male
}

```

Figure 1.4: Interface Definition

1.5 Scheme Directive

A directive enables to attach arbitrary information to any element. Directives are always placed behind the element they describe:

```

type Student implements User {
  id: Int!
  name: String! @defaultValue(value: "Iyadh")
  age: Int!
  gender: Gender!
  level: Int!
  major: String!
}

```

Figure 1.5: Applying directive to name

Chapter 2

Queries, Mutations and Subscriptions

2.1 Introduction

To manage the CRUD in different applications, GraphQL [\[Tea\]](#) offers different functions to manipulate these parts:

- query to fetch data
- mutation to add/update or delete data
- subscription to subscribe (listen) to a certain query

2.2 Queries

Queries in GraphQL are mainly about asking for specific fields on objects. It is completely flexible and lets the client decide what data is actually needed.

Queries enable different arguments to maintain their flexibility. Let's take an example of a user like we saw in the previous section (id, name and gender). We can do the following operations to get data:

2.2.1 Get all users

```
{
  user {
    name
  }
}
```

Figure 2.1: Get the names of all the users

2.2.2 Get user by id

```
{
  user(id: 1) {
    name
  }
}
```

Figure 2.2: Get name of the user of id 1

2.2.3 Get user by field name

```
{
  user(name : "iyadh") {
    age
  }
}
```

Figure 2.3: Get user ages with name "iyadh"

2.2.4 Get users with fragments

```
{
  user {
    ...userFields
  }
}

fragment userFields on User {
  name
  age
}
```

Figure 2.4: Get name and age of all users

2.2.5 Use variables

```
query getUsers($id: Int)
{
  user(id: $id) {
    name
  }
}

{
  "id" : 1
}
```

Figure 2.5: Set id as a variable that can be an input by the client

2.3 Mutations

Next to requesting information from a server, the majority of applications also need a way to make changes to the data that's currently stored in the backend. With GraphQL, these changes are made using so-called mutations. There generally are three kinds of mutations:

- creating new data
- updating existing data
- deleting existing data

Mutations follow the same syntactical structure as queries, but they always need to start with the mutation keyword. Here's an example for how we might create a new User:

```
mutation CreateNewUser($name: String!, $age: Int!) {  
  createUser(name: $name, age: $age) {  
    name  
    age  
    gender  
  }  
}
```

Figure 2.6: Creating a new user

2.4 Subscriptions

Another important requirement for many applications is to have a real-time connection to catch important events. For this use case, GraphQL offers the concept of subscriptions.

When a client subscribes to an event, it will initiate and hold a steady connection to the server. Whenever that particular event then actually happens, the server pushes the corresponding data to the client. Unlike queries and mutations that follow a typical “request-response-cycle”, subscriptions represent a stream of data sent over to the client.

Subscriptions are written using the same syntax as queries and mutations.

Chapter 3

GraphQL VS. REST

3.1 Introduction

GraphQL and REST [\[web\]](#) diverge in many ways. Thus, in this chapter, we will present the core difference between the two while explaining the concepts of each point.

3.2 Comparison Table

GraphQL	REST
Single Endpoint	Multiple Endpoints
GraphQL Scheme	URL Routes
No Overfetching or Underfetching	Overfetching and Underfetching
Fast Iterations on Front-End	Slow Iterations on Front-End
Fetches via HTTP GET Request via URL	Fetches via HTTP GET Request via URL
Returns JSON Data	Returns JSON Data

Table 3.1: GraphQL VS. REST

3.2.1 Overfetching

Overfetching implies that more information is accessed by a client than is currently required in the app. For instance, imagine a screen that needs to show a list of users with only their names. This app will usually reach the endpoint of `/users` in a REST API and receive a JSON array of user data. However this response can provide more information about the users that are returned, such as their birthdays or addresses - information that is irrelevant to the consumer because it only needs to show the names of the users.

3.2.2 Underfetching

Underfetching typically implies that not enough of the information needed is supplied by a particular endpoint. To fetch all it wants, the customer will have to make more requests. This can lead to a situation where a client wants to download a list of components first but then needs to make one additional request per component to retrieve the necessary information. As an example consider that the last three followers per user will also need to be displayed in the same app. The additional endpoint `/users/followers` is provided by the API. The app would have to make one request to the `/users` endpoint to be able to view the appropriate data and then reach the `/users/followers` endpoint for each user.

3.2.3 Single Endpoint VS. Multiple Endpoints

You can usually collect the data with a REST API by accessing several endpoints. In this case, the endpoint could be `/users/<id>` to fetch the initial user data. Second, there is likely to be an endpoint of `/users/<id>/posts` that returns all posts to a user. The third endpoint will then be the `/users/<id>/followers` that returns a list of followers per user.




Figure 3.1: REST Example Endpoint



Figure 3.2: REST Example Endpoint 2

On the other hand, in GraphQL, you simply send a single query that contains the basic data specifications to the GraphQL server. The server then responds, when these requirements are met, with a JSON object.




Figure 3.3: GraphQL Example Endpoint

Chapter 4

Practical case: Hasura

4.1 Introduction

In this section, we are going to present multiple steps to configure Hasura. Although there are different ways to do so, we chose to use Docker in order to make the installation smoother and easier. [?]

4.2 Steps

4.2.1 Create a new docker-compose file locally

```
(base) MacBook-Pro-de-macbook-3:hasura_dev macbook$ touch docker-compose.yml
```

Figure 4.1: File creating command

4.2.2 Copy the docker-compose content from GraphQL Github Repository

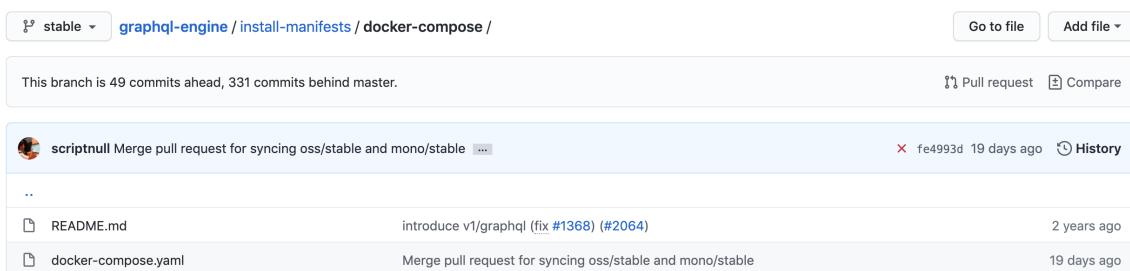
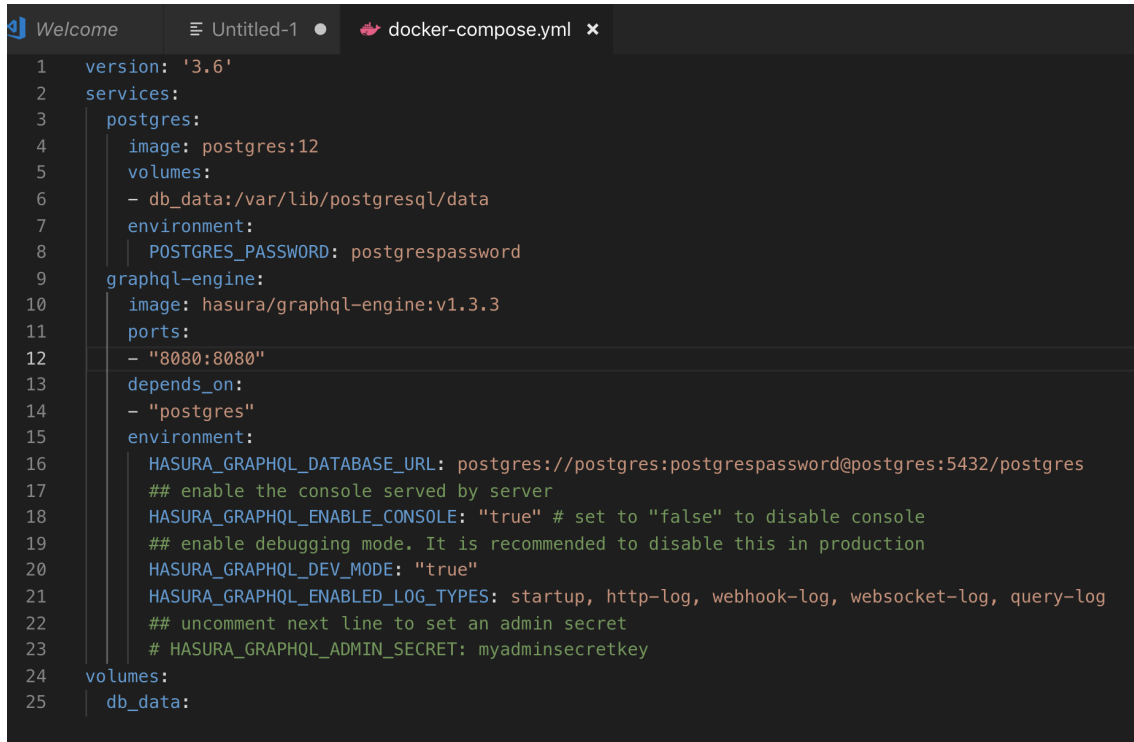


Figure 4.2: Github File

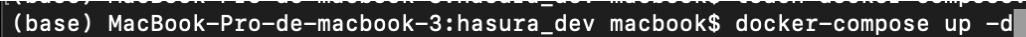
4.2.3 Check the local file variables

A screenshot of a code editor window with a dark theme. The window has three tabs: 'Welcome', 'Untitled-1', and 'docker-compose.yml'. The 'docker-compose.yml' tab is active, showing a YAML configuration for two services: 'postgres' and 'graphql-engine'. The 'postgres' service uses the 'postgres:12' image and has a volume 'db_data' mapped to '/var/lib/postgresql/data'. The 'graphql-engine' service uses the 'hasura/graphql-engine:v1.3.3' image, has port '8080' mapped to '8080', depends on the 'postgres' service, and has several environment variables including 'HASURA_GRAPHQL_DATABASE_URL', 'HASURA_GRAPHQL_ENABLE_CONSOLE', 'HASURA_GRAPHQL_DEV_MODE', 'HASURA_GRAPHQL_ENABLED_LOG_TYPES', and 'HASURA_GRAPHQL_ADMIN_SECRET'. There is also a 'volumes' section at the bottom with 'db_data' defined.

```
1 version: '3.6'
2 services:
3   postgres:
4     image: postgres:12
5     volumes:
6     - db_data:/var/lib/postgresql/data
7     environment:
8       POSTGRES_PASSWORD: postgrespassword
9   graphql-engine:
10    image: hasura/graphql-engine:v1.3.3
11    ports:
12    - "8080:8080"
13    depends_on:
14    - "postgres"
15    environment:
16      HASURA_GRAPHQL_DATABASE_URL: postgres://postgres:postgrespassword@postgres:5432/postgres
17      ## enable the console served by server
18      HASURA_GRAPHQL_ENABLE_CONSOLE: "true" # set to "false" to disable console
19      ## enable debugging mode. It is recommended to disable this in production
20      HASURA_GRAPHQL_DEV_MODE: "true"
21      HASURA_GRAPHQL_ENABLED_LOG_TYPES: startup, http-log, webhook-log, websocket-log, query-log
22      ## uncomment next line to set an admin secret
23      # HASURA_GRAPHQL_ADMIN_SECRET: myadminsecretkey
24 volumes:
25   db_data:
```

Figure 4.3: Docker-compose local file

4.2.4 Run the docker-compose up command

A screenshot of a terminal window. The prompt is '(base) MacBook-Pro-de-macbook-3:hasura_dev macbook\$'. The command 'docker-compose up -d' has been entered, and the cursor is at the end of the line.

```
(base) MacBook-Pro-de-macbook-3:hasura_dev macbook$ docker-compose up -d
```

Figure 4.4: Command Running

4.2.5 Open the console on port 8080

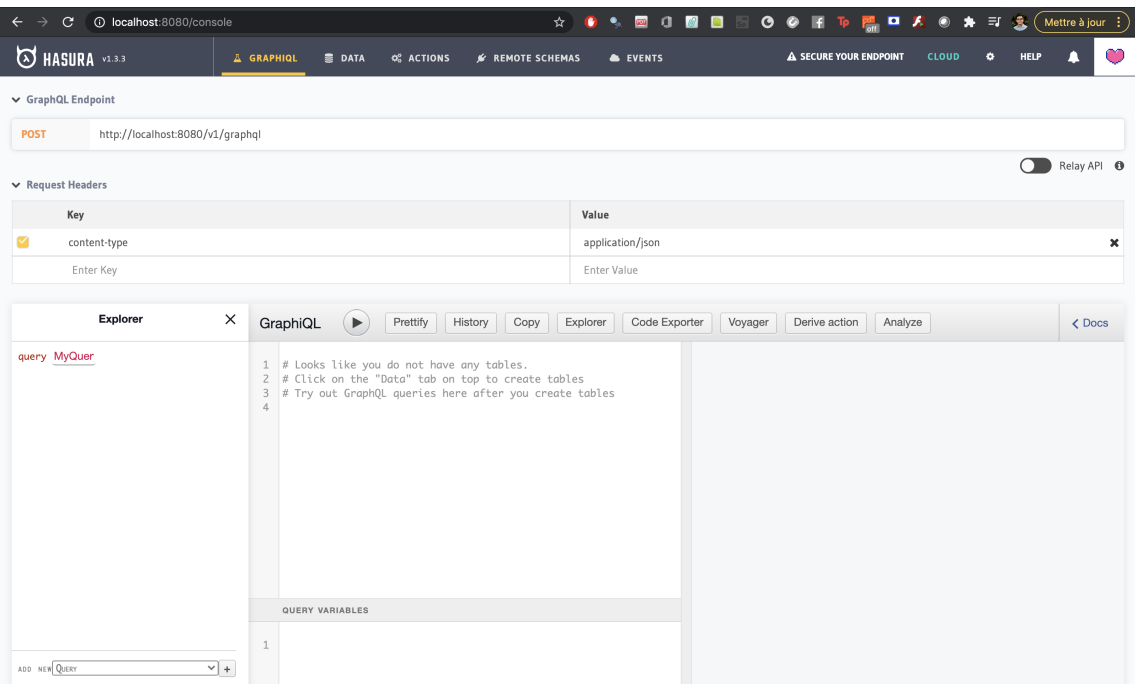


Figure 4.5: GraphQL Console

4.2.6 Create a new table: Users

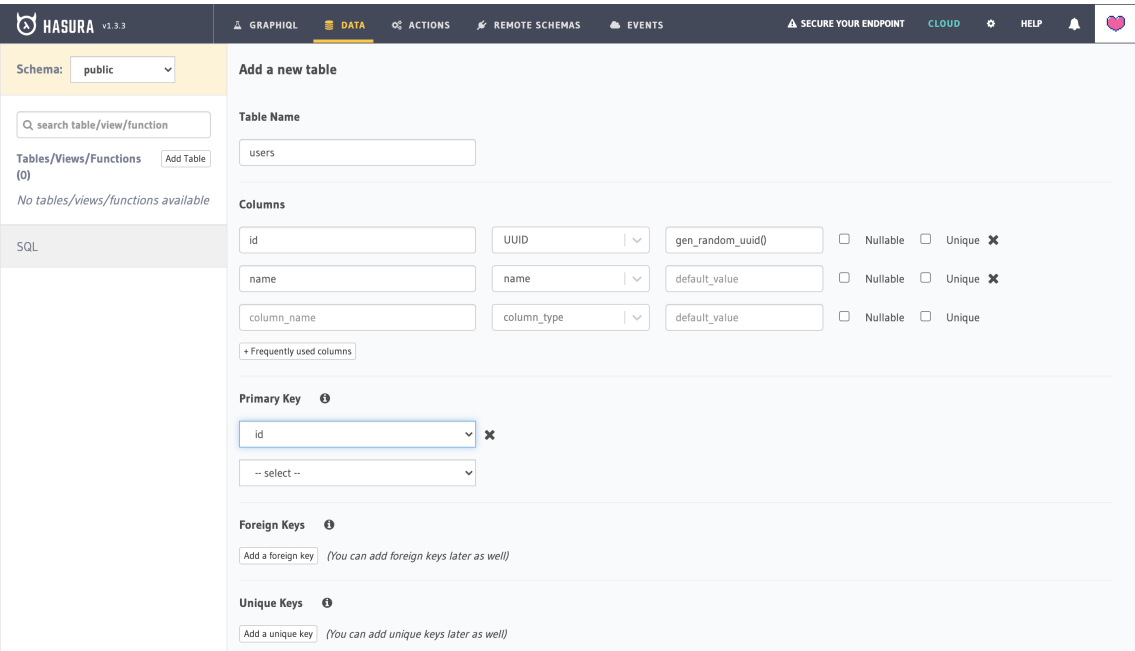


Figure 4.6: Users Table Creation

4.2.7 Create a new table: Courses

The screenshot shows the Hasura v1.3.3 interface. On the left, there's a sidebar with 'Schema: public' and a search bar. Below it, 'Tables/Views/Functions (1)' lists 'users'. The main area is titled 'Add a new table'. Under 'Table Name', 'courses' is entered. The 'Columns' section lists five columns: 'id' (UUID, gen_random_uuid(), Nullable, Unique), 'title' (name, default_value, Nullable, Unique), 'summary' (Text, default_value, Nullable, Unique), 'author_id' (UUID, example: gen_random_uuid(), Nullable, Unique), and 'column_name' (column_type, default_value, Nullable, Unique). Each column has checkboxes for 'Nullable' and 'Unique'.

Figure 4.7: Courses Table Creation

4.2.8 Configure Primary and Foreign Keys in Courses

The screenshot shows the 'Primary Key' and 'Foreign Keys' configuration for the 'courses' table. The 'Primary Key' is set to 'id'. The 'Foreign Keys' section shows a configuration for 'author_id' referencing 'users.id' with 'restrict' on update and delete violations. The 'Reference Schema' is 'public' and the 'Reference Table' is 'users'. The 'From' column is 'author_id' and the 'To' column is 'id'. The 'On Update Violation' and 'On Delete Violation' are both set to 'restrict'.

Figure 4.8: Key configuration

4.2.9 Define the relationship between tables

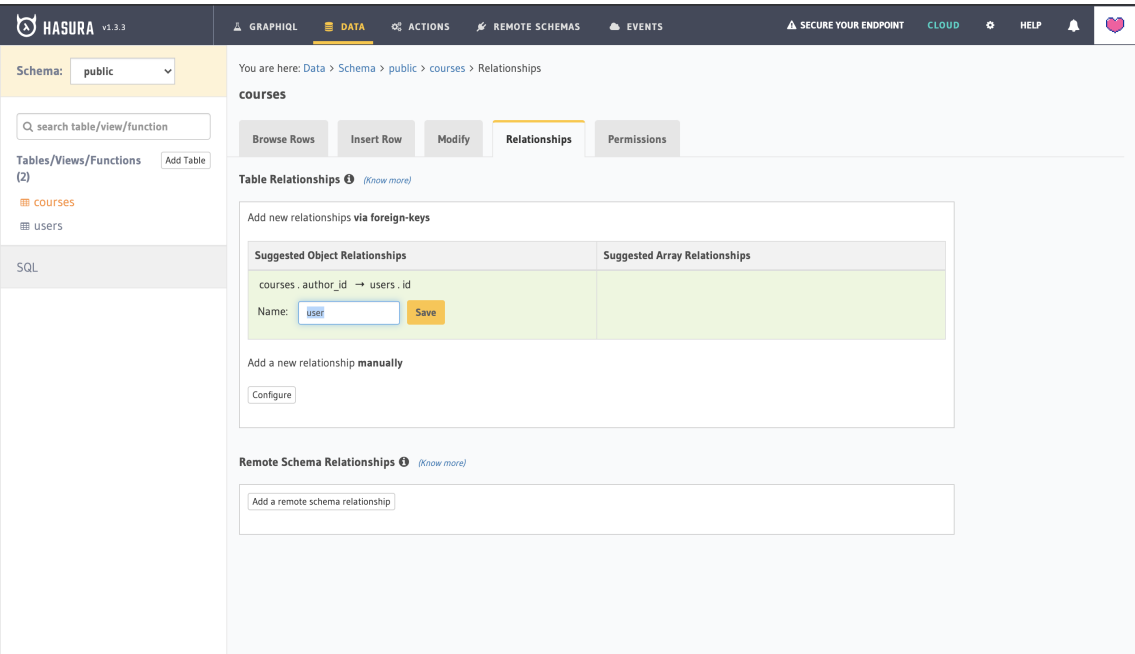


Figure 4.9: Relationship Definition

4.2.10 Mutation: Adding new users



Figure 4.10: Adding users

4.2.11 Mutation: Adding new courses

```
mutation createCourse {
  insert_courses(
    objects:[{title: "E-Services",
              summary:"This is an E-Services Course",
              author_id: "1eaf121f-d142-4a4a-b878-7569c6b81796"},
             {title: "Big Data",
              summary:"This is a Big Data Course",
              author_id: "1eaf121f-d142-4a4a-b878-7569c6b81796"}]
  ) {
    returning {
      id
      user {
        name
        occupation
      }
    }
  }
}
```

```
{
  "data": {
    "insert_courses": {
      "returning": [
        {
          "id": "5a09df3f-a075-460c-af99-599e0e685175",
          "user": {
            "name": "Lilia Sfaxi",
            "occupation": "Professor"
          }
        },
        {
          "id": "20c76a12-7ecd-4544-b285-3afc85b7346a",
          "user": {
            "name": "Lilia Sfaxi",
            "occupation": "Professor"
          }
        }
      ]
    }
  }
}
```

Figure 4.11: Adding courses

4.2.12 Subscription: Subscribing to users

The api-explorer keeps on listening for new users and appends them to the response JSON in real-time.

```
subscription sub{
  users{
    id
    name
  }
}
```

```
{
  "data": {
    "users": [
      {
        "id": "7f888ba9-15d0-4e08-8e99-37b27098fbf8",
        "name": "Iyadh Khalfallah"
      },
      {
        "id": "3c32b1fa-7ac6-4f3f-905d-229e7e636cc7",
        "name": "Iyadh Khalfallah"
      },
      {
        "id": "1eaf121f-d142-4a4a-b878-7569c6b81796",
        "name": "Lilia Sfaxi"
      }
    ]
  }
}
```

Figure 4.12: User subscription

Conclusion and perspectives

GraphQL is an increasingly important technology. We can see its magic through different pros compared to REST, from a single endpoint to communicate, to the avoidance on Overfetching and Underfetching. In some papers, it is shown that GraphQL can reduce the size of the JSON documents returned by REST APIs in 94% (in number of fields) and in 99% (in number of bytes) [GB19]. Nevertheless, GraphQL has numerous drawbacks that can harm the backend system in different ways. Thus, before choosing the backend technology, one might study different aspects of his system.

Bibliography

- [All18] Alligator.io. A graphql sdl reference. digitalocean.com/community/tutorials/graphql-graphql-sdl, 2018.
- [Bur17] Nikolas Burk. GraphQL sdl — schema definition language. prisma.io/blog/graphql-sdl-schema-definition-language-6755bcb9ce51, 2017.
- [GB19] Marco Tulio Valente ASERG Group Department of Computer Science Federal University of Minas Gerais Brazil Gleison Brito, Thais Mombach. Migrating to graphql: A practical assessment. 2019.
- [Tea] GraphQL Team. Introduction to graphql. graphql.org/learn/.
- [web]